From Claude to Tax

How AI evolves in three stages — what we can already see in software, and what we're building at Meta to bring it to tax.

A practical look at the infrastructure that makes it real.

Coding is the leading indicator

Claude Code made AI-native development the default. Coding is where AI hit first — so it's the map for what's coming next, including tax.

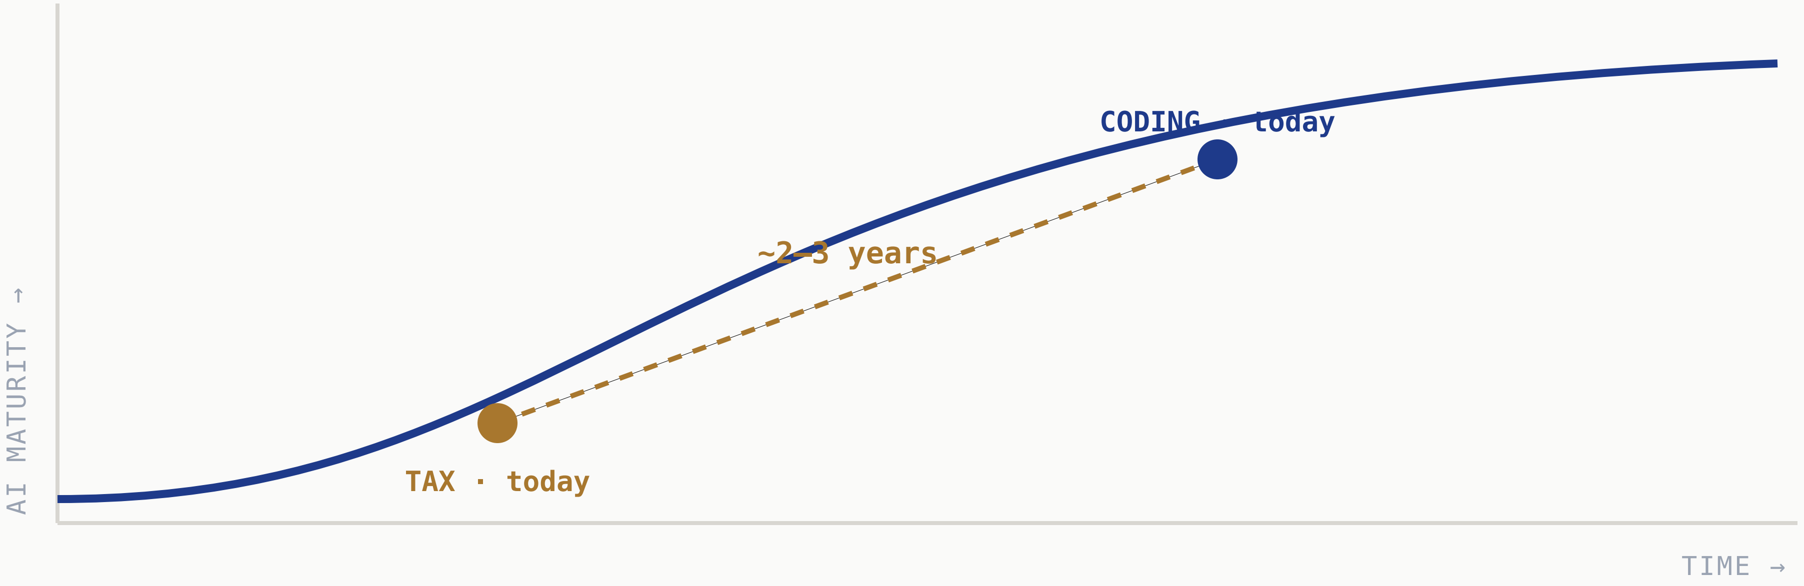

Same curve, a few years apart — **watch coding and you can read tax's map ahead.**

THE PATTERN

AI evolves in three stages — climbed in order

01

AI-Native Foundation

AI & data built in, not bolted on. Copilots assist; humans own.

02

Continuous Learning

Learns from feedback & outcomes, and adapts on its own. Accuracy compounds.

03

Proactive Management

Acts *before* you ask. Near-autonomous — but a human still approves.

↑ EACH STEP NEEDS THE ONE BELOW

You can't skip a step. The order isn't optional.

WHERE CODING STANDS

Score the three stages, 0–100

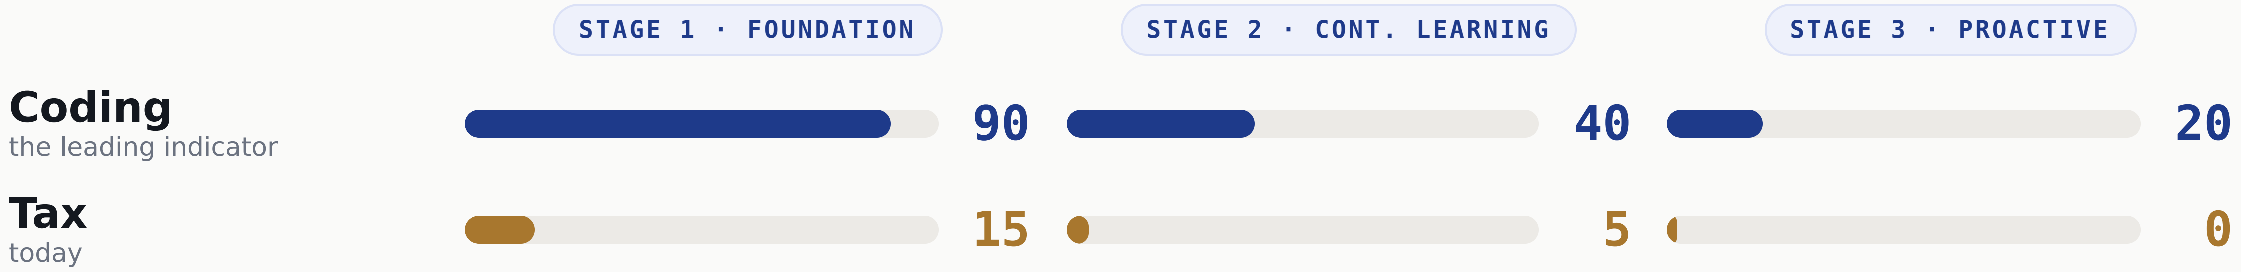

Coding has finished Stage 1 and is mid-Stage 2. Tax is at the very start —
same curve, ~2–3 years apart.

WHY THIS MATTERS FOR TAX

Coding figures it out first — then it diffuses

THE FRONTIER

Coding

Most investment, most talent — the harness gets invented here first.

DIFFUSES →

- **Tax** adopts the proven harness as it arrives
- **Finance & Legal** next in line
- **Other verticals** follow the same path

The harness for Stage 2 & 3 doesn't fully exist yet — anywhere. **We don't need to invent it. We need to be ready when it arrives.**

OUR PLAY AT META

So we start at Stage 1 — across the whole tax lifecycle

We're building **one** AI-native foundation — and every stage of the tax lifecycle draws from it.

Tax Planning

Scenario modeling & structuring, grounded in our positions and the law.

Compliance & Reporting

High-volume, rules-based work on real-time data — the natural first mover.

Tax Audit

Controversy & defense on the same lineage, sources, and document trail.

1 **One AI-Native Foundation** shared data · grounding · governance — under all three

Foundation first. Build it once, build it well.

WHY FOUNDATION-FIRST PAYS OFF

Connect end-to-end — a flywheel kicks in

Connect Planning → Compliance → Audit on one foundation and each loop compounds — **audit closes the loop.**

1 · Connect end-to-end
Planning → Compliance → Audit on one foundation

4 · Foundation compounds
Cleaner data → better planning & compliance → more to connect

↻
FLYWHEEL
audit closes the loop

2 · Audit-ready lineage
Every number traces to source — defensible on demand

3 · Feedback captured
Audit outcomes & corrections become learning signal

Each turn makes the data cleaner and the foundation smarter — **the exact feedback engine Stage 2 will run on.**

UNDER THE HOOD

Every layer has an industry — and a Meta system

The "AI" is just the models. Everything else is infrastructure — and we already run an internal version of nearly all of it.

LAYER	INDUSTRY	META · INTERNAL
UI / experience	web & app · copilot UI · workflow	React · Relay
Models (the AI)	Claude · GPT · Gemini	Llama · PyTorch
Retrieval / vector	Pinecone · Weaviate · pgvector	FAISS
Graph / knowledge graph	Neo4j · TigerGraph · Neptune	TAO · Dragon
Master data & catalog	Collibra · Alation · Atlan	Nemo (catalog) · lineage
Integration & connectors	SAP/Oracle · Vertex/Avalara · MuleSoft	internal finance systems
Streaming & pipelines	Kafka/Confluent · Flink · Fivetran	Scribe · Dataswarm · Scuba
Data platform / lakehouse	Snowflake · Databricks · BigQuery	Presto · Hive · Tectonic
Governance & security	Immuta · OneTrust · Okta	Privacy Aware Infrastructure
Eval & observability	LangSmith · Arize · Braintrust	Scuba · ODS

The moat isn't the model — it's already running **all ten layers.**

MAKE IT CONCRETE

Why the model alone isn't enough

“Did we charge the right sales tax on this $2M order shipped to Texas?”

The model alone

Knows Texas sales tax *in general*. But not **your** order, your nexus, today's rate — or how to prove it.

With the foundation, it can…

Data platform	pull the real order — $2M, ship-to TX, the product SKU
Graph	which legal entity sold it · TX state→county→city · is the product taxable
Retrieval	today's Texas rule & rate — **with a citation**
Models	reason over all of it to an answer
Governance	check who may see it · log every step
Eval + lineage	confidence check; unsure → human · every number traces back

The model is one slice of the answer. The **infrastructure is the rest** — and the part that holds up in an audit.

THE HONEST PART

Hard for most — but we have a head start

Why it's hard

- Fragmented, messy data across ERPs & jurisdictions
- Real-time / streaming is a specialized skill set
- Talent fluent in tax *and* AI is doubly rare
- Split ownership: Tax vs. Finance Systems vs. IT
- Tax is a cost center — hard to fund infra

Why Meta is positioned

- Hyperscale data infrastructure already exists
- Real-time / event-driven by default
- World-class platform, data & ML engineers
- AI platform & compute in-house, not procured
- Engineering-owned tax tech — ships like a product team

The barriers most teams hit are largely **already solved here** — so we get to start where others can't.

WHY META LEADS THIS

The transformation takes all three at once

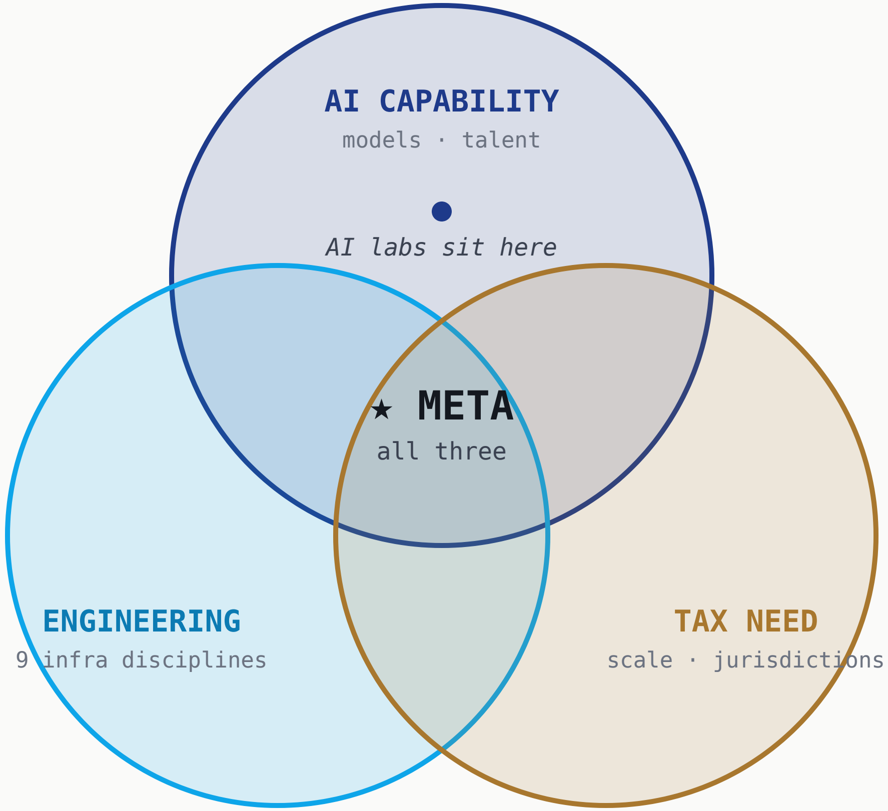

Why not the AI labs?

- Anthropic & OpenAI have world-class AI — but that's **one circle, not three**.
- Lighter **broad-engineering** investment — focused on model R&D, not a full enterprise data platform.
- Early, **small tax footprint** — one is just standing up a tax engine (Vertex).

You can't model your way there. It takes **AI × engineering × real tax need — together.**

THE TAKEAWAY

Build the foundation now. Be ready when the harness arrives.

Coding shows the path. The three stages are real. The winning move in tax is Stage 1 — done well, across the lifecycle.

taxbenchmark.ai

Drop a qr.png here before the talk.

Benchmark your tax-tech function

taxbenchmark.ai

Thank you · contact@taxbenchmark.ai